

# Chapter 3: Processes

---





# Chapter 3: Processes

---

- n Process Concept
  - | Process states, process control block, scheduling queues
- n Schedulers
- n Context switch
- n Process creation
- n Process Termination






# Process Concept

---

- n An operating system executes a variety of programs:
  - | Batch system – **jobs**
  - | Time-shared systems – **user programs** or **tasks**
- n Textbook uses the terms **job** and **process** almost interchangeably
- n **Process** – a program in execution; process execution must progress in sequential fashion
- n Multiple parts
  - | The program code, also called **text section**
  - | Current activity including **program counter**, processor registers
  - | **Stack** containing temporary data
    - Function parameters, return addresses, local variables
  - | **Data section** containing global variables
  - | **Heap** containing memory dynamically allocated during run time





# Process Concept (Cont.)


---

- n Program is **passive** entity stored on disk (**executable file**), process is **active**
  - | Program becomes process when executable file loaded into memory
- n Execution of program started via GUI mouse clicks, command line entry of its name, etc
- n One program can be several processes
  - | Consider multiple users executing the same program





# Process in Memory





# Process State

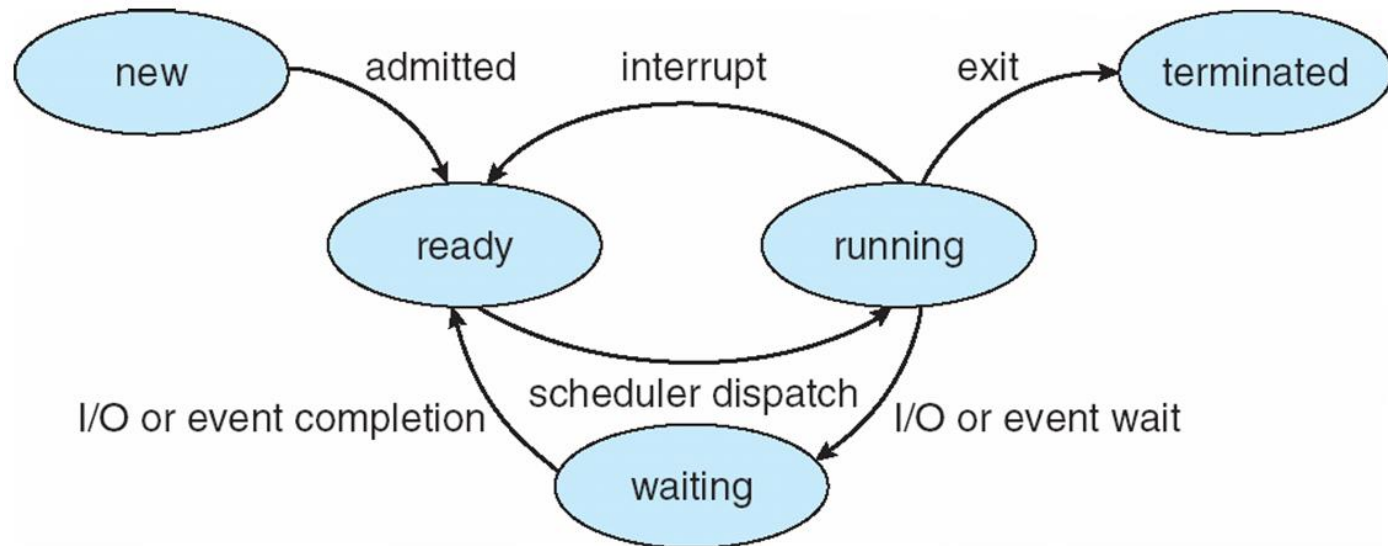
---

- n As a process executes, it changes **state**
  - | **new**: The process is being created
  - | **running**: Instructions are being executed
  - | **waiting**: The process is waiting for some event to occur
  - | **ready**: The process is waiting to be assigned to a processor
  - | **terminated**: The process has finished execution





# Diagram of Process State






# Process Control Block (PCB)

Information associated with each process  
(also called **task control block**)


- n Process state – running, waiting, etc
- n Program counter – location of instruction to next execute
- n CPU registers – contents of all process-centric registers
- n CPU scheduling information- priorities, scheduling queue pointers
- n Memory-management information – memory allocated to the process
- n Accounting information – CPU used, clock time elapsed since start, time limits
- n I/O status information – I/O devices allocated to process Example: list of open files







# CPU Switch From Process to Process





# Threads

---

- n So far, process has a single thread of execution
- n Consider having multiple program counters per process
  - | Multiple locations can execute at once
    - ▶ Multiple threads of control -> **threads**
- n Must then have storage for thread details, multiple program counters in PCB
- n See next chapter





# Process Scheduling


---

- n Maximize CPU use, quickly switch processes onto CPU for time sharing
- n **Process scheduler** selects among available processes for next execution on CPU
- n Maintains **scheduling queues** of processes
  - | **Job queue** – set of all processes in the system
  - | **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
  - | **Device queues** – set of processes waiting for an I/O device
  - | Processes migrate among the various queues






# Ready Queue And Various I/O Device Queues





# Representation of Process Scheduling

n **Queueing diagram** represents queues, resources, flows





# Schedulers


- n **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU
  - | Sometimes the only scheduler in a system
  - | Short-term scheduler is invoked frequently (milliseconds)  $\Rightarrow$  (must be fast)
- n **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
  - | Long-term scheduler is invoked infrequently (seconds, minutes)  $\Rightarrow$  (may be slow)
  - | The long-term scheduler controls the **degree of multiprogramming** (the number of processes in memory)
- n Processes can be described as either:
  - | **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
  - | **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- n Long-term scheduler strives for good ***process mix***





# Addition of Medium Term Scheduling

- n **Medium-term scheduler** can be added if degree of multiple programming needs to decrease
  - I Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**





# Context Switch

- n When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- n **Context** of a process represented in the PCB
- n Context-switch time is overhead; the system does no useful work while switching
  - | The more complex the OS and the PCB → the longer the context switch
- n Time dependent on hardware support
  - | Some hardware provides multiple sets of registers per CPU → multiple contexts loaded at once







# Operations on Processes

---

System must provide mechanisms for:  
process creation,  
process termination,






# Process Creation

- n **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- n Generally, process identified and managed via a **process identifier (pid)**
- n Resource sharing options
  - | Parent and children share all resources
  - | Children share subset of parent's resources
  - | Parent and child share no resources
- n Execution options
  - | Parent and children execute concurrently
  - | Parent waits until children terminate






# A Tree of Processes in Linux





# Process Creation (Cont.)

- n Address space
  - | Child duplicate of parent
  - | Child has a program loaded into it
- n UNIX examples
  - | **fork()** system call creates new process
  - | **exec()** system call used after a **fork()** to replace the process' memory space with a new program





# Process Termination

---

- n Process executes last statement and then asks the operating system to delete it using the **exit()** system call.
  - | Returns status data from child to parent (via **wait()**)
  - | Process' resources are deallocated by operating system
- n Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:
  - | Child has exceeded allocated resources
  - | Task assigned to child is no longer required
  - | The parent is exiting and the operating systems does not allow a child to continue if its parent terminates





# Process Termination

- n Some operating systems do not allow child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.
  - | **cascading termination.** All children, grandchildren, etc. are terminated.
  - | The termination is initiated by the operating system.
- n The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the pid of the terminated process  

```
pid = wait(&status);
```
- n If no parent waiting (did not invoke `wait()`) process is a **zombie**
- n If parent terminated without invoking `wait`, process is an **orphan**





# CODE 1/3

---

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
int main(){
    // Fork a child process
    pid_t child_pid = fork();
    int status;
    if (child_pid == -1) {
        // Fork failed
        printf("Fork failed");
        exit(EXIT_FAILURE);
    }
```





## CODE 2/3

---

```
if (child_pid > 0) {  
    // Parent process  
    printf("Parent process (PID: %d)\n", getpid());  
    printf("Child process created with PID: %d\n", child_pid);  
    wait(&status);  
    printf("Child exited with status: %d\n", status);  
}  
  
else if (child_pid == 0) {  
    // Child process  
    printf("Child process (PID: %d)\n", getpid());  
    exit(EXIT_FAILURE);  
    //exit(EXIT_SUCCESS);  
}
```







## CODE 3/3

---

```
// Code executed by both parent and child processes
printf("This code is executed by both parent and child processes.\n");

return 0;
}
```





# OUTPUT

---

Child process (PID: 23787)

Parent process (PID: 23786)

Child process created with PID: 23787

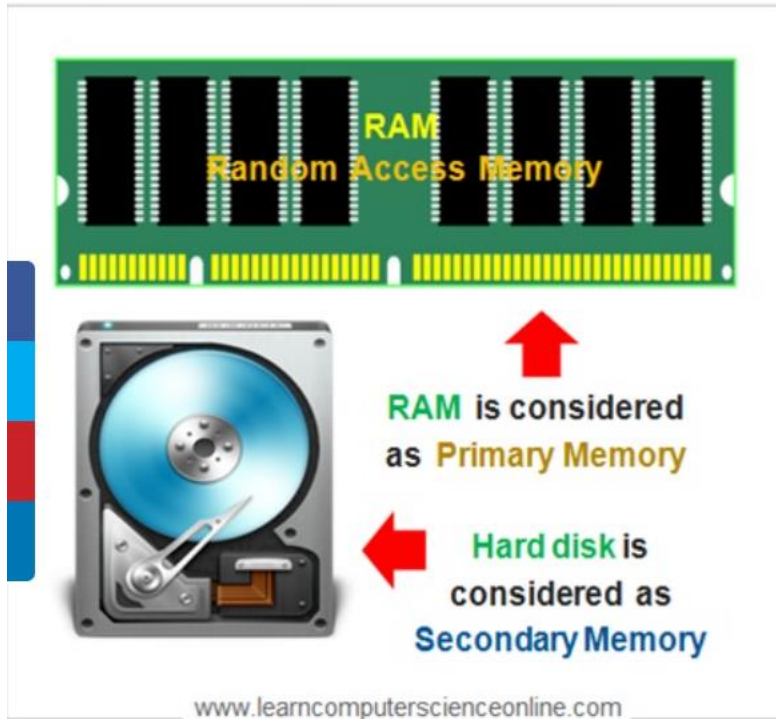
Child exited with status: 256

This code is executed by both parent and child processes.





## Computer Primary And Secondary Memory



## Computer Memory Hierarchy

